

Neural Network Performance Analysis

This notebook explores various neural network architectures to predict sentiment in the IMDB dataset. Different models were experimented with, varying in hidden layers, hidden units, and loss functions. Below is the analysis of these models.

```
In [41]: %matplotlib inline
!jupyter nbconvert --to html Neural_Network_Assignment.ipynb
from google.colab import files
files.download('Neural_Network_Assignment.html')
```

```
[NbConvertApp] Converting notebook Neural_Network_Assignment.ipynb to html
[NbConvertApp] Writing 593635 bytes to Neural_Network_Assignment.html
```

Model with One Hidden Layer

- **Overview:** The model was trained using a single hidden layer with 32 units and the `tanh` activation function.
- **Performance:** The training accuracy slowly improved but remained unstable across the epochs. Validation accuracy fluctuated and settled around 50%, indicating possible underfitting.
- **Conclusion:** Using only one hidden layer resulted in suboptimal performance, which might be due to insufficient model complexity.

```
In [ ]: # Import necessary modules
from tensorflow.keras import models, layers, regularizers
from tensorflow.keras.datasets import imdb
from tensorflow.keras.preprocessing.sequence import pad_sequences

# Load the IMDB dataset
(train_data, train_labels), (test_data, test_labels) = imdb.load_data(num_words=10000)

# Preprocess the data (pad sequences to make them the same length)
partial_x_train = pad_sequences(train_data[:20000], maxlen=256)
partial_y_train = train_labels[:20000]
x_val = pad_sequences(train_data[20000:], maxlen=256)
y_val = train_labels[20000:]

# One hidden layer model
model_one_hidden_layer = models.Sequential([
    layers.Dense(32, activation="tanh", kernel_regularizer=regularizers.l2(0.001)),
    layers.Dense(1, activation="sigmoid")
])

# Compile the model
model_one_hidden_layer.compile(optimizer="rmsprop", loss="binary_crossentropy", metrics=["accuracy"])

# Train the model
history_one_hidden_layer = model_one_hidden_layer.fit(partial_x_train, partial_y_train, epochs=20, batch_size=5)

import matplotlib.pyplot as plt

def plot_history(history, title):
    history_dict = history.history
    loss_values = history_dict["loss"]
    val_loss_values = history_dict["val_loss"]
    acc = history_dict.get("accuracy", None)
    val_acc = history_dict.get("val_accuracy", None)
    epochs = range(1, len(loss_values) + 1)

    # Plot training and validation loss
    plt.plot(epochs, loss_values, "bo", label="Training loss")
    plt.plot(epochs, val_loss_values, "b", label="Validation loss")
    plt.title(f"Training and validation loss: {title}")
    plt.xlabel("Epochs")
    plt.ylabel("Loss")
    plt.legend()
    plt.show()

    # If accuracy is available, plot training and validation accuracy
    if acc and val_acc:
        plt.clf()
        plt.plot(epochs, acc, "bo", label="Training accuracy")
        plt.plot(epochs, val_acc, "b", label="Validation accuracy")
        plt.title(f"Training and validation accuracy: {title}")
        plt.xlabel("Epochs")
        plt.ylabel("Accuracy")
        plt.legend()
        plt.show()

# Plot the results for one hidden layer
plot_history(history_one_hidden_layer, "One Hidden Layer")
```

Epoch 1/20
40/40 ————— 2s 15ms/step - accuracy: 0.5035 - loss: 1.0487 - val_accuracy: 0.5014 - val_loss: 0.9540

Epoch 2/20
40/40 ————— 0s 4ms/step - accuracy: 0.4985 - loss: 0.9489 - val_accuracy: 0.5062 - val_loss: 0.9003

Epoch 3/20
40/40 ————— 0s 4ms/step - accuracy: 0.5090 - loss: 0.8887 - val_accuracy: 0.5110 - val_loss: 0.8582

Epoch 4/20
40/40 ————— 0s 6ms/step - accuracy: 0.5113 - loss: 0.8534 - val_accuracy: 0.5012 - val_loss: 0.8324

Epoch 5/20
40/40 ————— 0s 4ms/step - accuracy: 0.5071 - loss: 0.8221 - val_accuracy: 0.5036 - val_loss: 0.8053

Epoch 6/20
40/40 ————— 0s 4ms/step - accuracy: 0.5102 - loss: 0.8016 - val_accuracy: 0.5050 - val_loss: 0.7877

Epoch 7/20
40/40 ————— 0s 4ms/step - accuracy: 0.5132 - loss: 0.7800 - val_accuracy: 0.5054 - val_loss: 0.7722

Epoch 8/20
40/40 ————— 0s 4ms/step - accuracy: 0.5107 - loss: 0.7690 - val_accuracy: 0.4950 - val_loss: 0.7651

Epoch 9/20
40/40 ————— 0s 4ms/step - accuracy: 0.5102 - loss: 0.7590 - val_accuracy: 0.4980 - val_loss: 0.7540

Epoch 10/20
40/40 ————— 0s 4ms/step - accuracy: 0.5131 - loss: 0.7505 - val_accuracy: 0.5004 - val_loss: 0.7483

Epoch 11/20
40/40 ————— 0s 5ms/step - accuracy: 0.5188 - loss: 0.7445 - val_accuracy: 0.5012 - val_loss: 0.7438

Epoch 12/20
40/40 ————— 0s 4ms/step - accuracy: 0.5203 - loss: 0.7405 - val_accuracy: 0.4944 - val_loss: 0.7392

Epoch 13/20
40/40 ————— 0s 4ms/step - accuracy: 0.5158 - loss: 0.7374 - val_accuracy: 0.5164 - val_loss: 0.7367

Epoch 14/20
40/40 ————— 0s 6ms/step - accuracy: 0.5174 - loss: 0.7336 - val_accuracy: 0.5108 - val_loss: 0.7331

Epoch 15/20
40/40 ————— 0s 4ms/step - accuracy: 0.5156 - loss: 0.7311 - val_accuracy: 0.5118 - val_loss: 0.7307

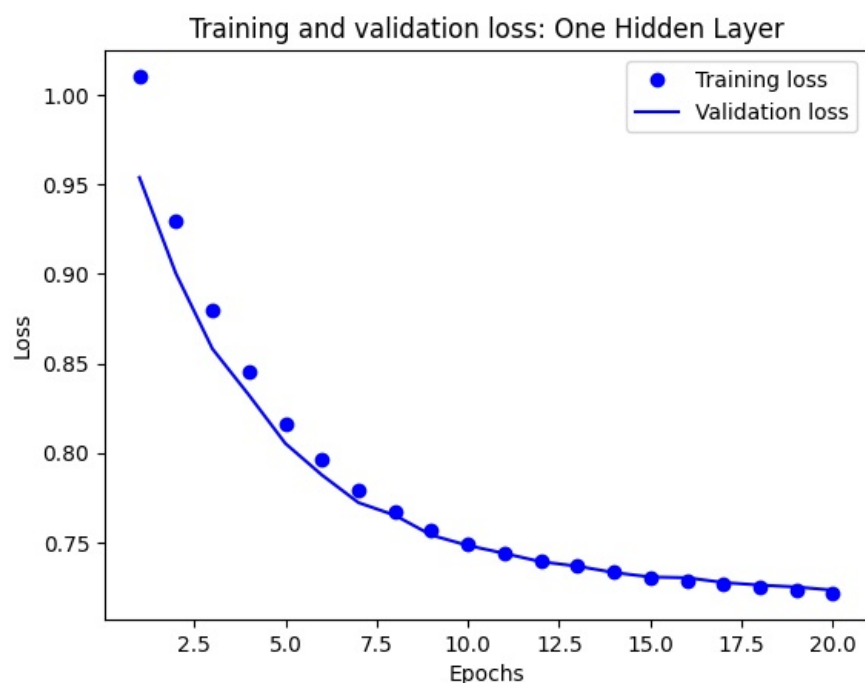
Epoch 16/20
40/40 ————— 0s 4ms/step - accuracy: 0.5213 - loss: 0.7297 - val_accuracy: 0.5120 - val_loss: 0.7302

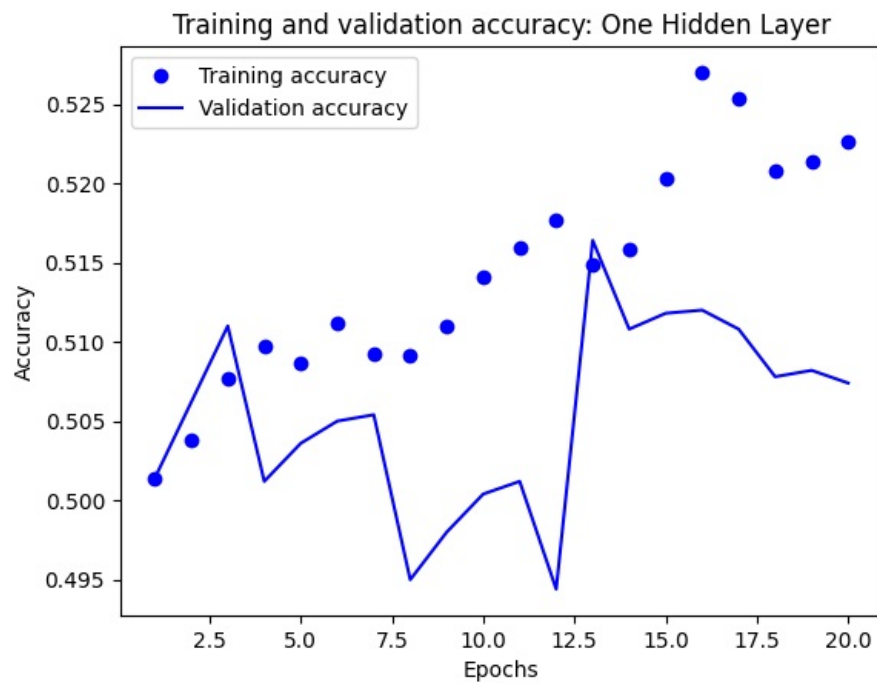
Epoch 17/20
40/40 ————— 0s 4ms/step - accuracy: 0.5219 - loss: 0.7273 - val_accuracy: 0.5108 - val_loss: 0.7276

Epoch 18/20
40/40 ————— 0s 5ms/step - accuracy: 0.5261 - loss: 0.7244 - val_accuracy: 0.5078 - val_loss: 0.7261

Epoch 19/20
40/40 ————— 0s 4ms/step - accuracy: 0.5239 - loss: 0.7232 - val_accuracy: 0.5082 - val_loss: 0.7251

Epoch 20/20
40/40 ————— 0s 4ms/step - accuracy: 0.5207 - loss: 0.7225 - val_accuracy: 0.5074 - val_loss: 0.7233





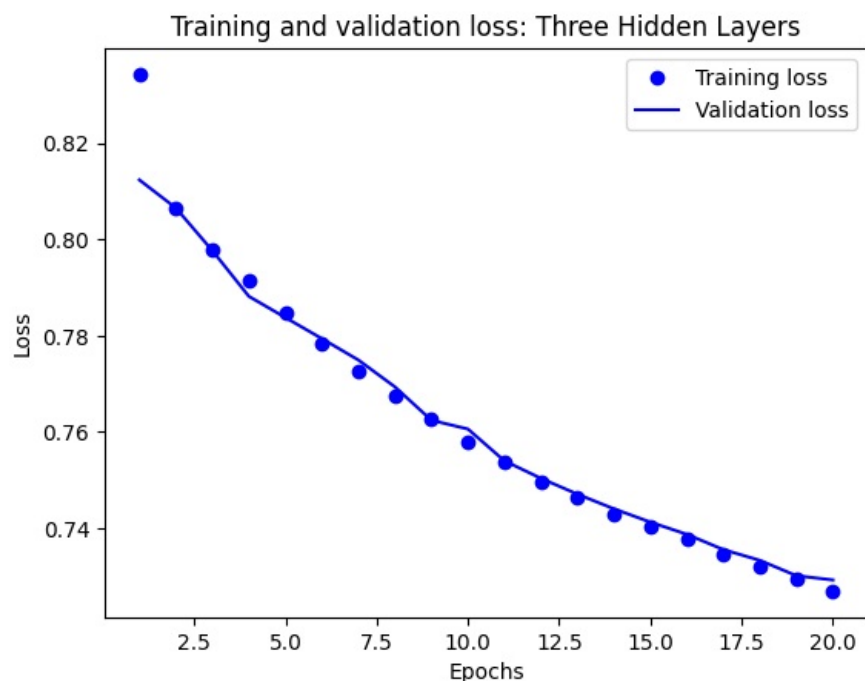
Model with Three Hidden Layers

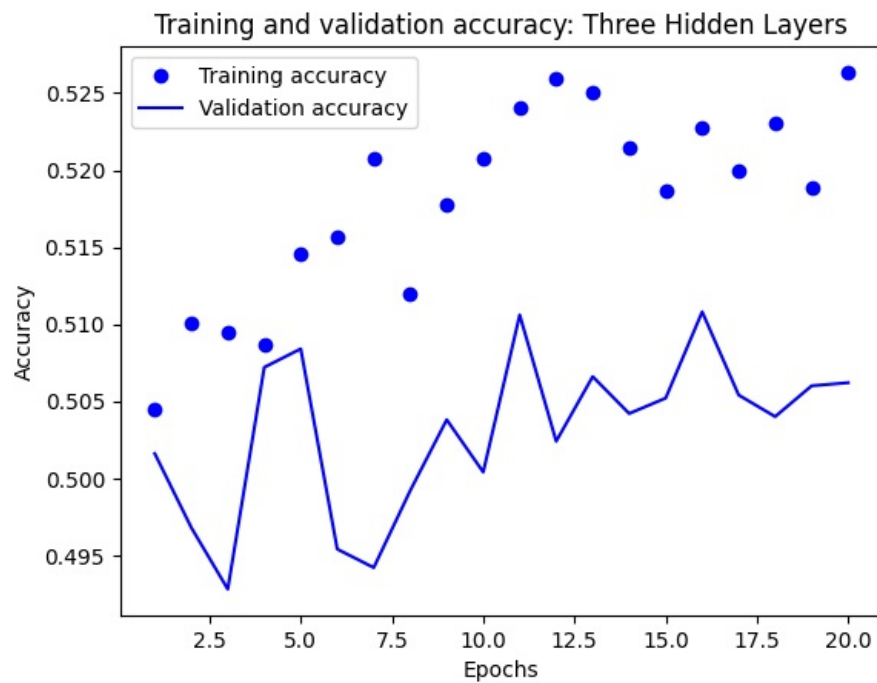
- **Overview:** This model had three hidden layers with 32 units each.
- **Performance:** Training accuracy was slightly better than the one-hidden-layer model, but the validation accuracy hovered around 50%. Training loss consistently decreased, but validation loss did not improve much.
- **Conclusion:** Increasing the model complexity did not lead to better generalization, likely due to overfitting or a lack of sufficient regularization.

```
In [ ]: # Three hidden layers model
model_three_hidden_layers = models.Sequential([
    layers.Dense(32, activation="tanh", kernel_regularizer=regularizers.l2(0.001)),
    layers.Dense(32, activation="tanh", kernel_regularizer=regularizers.l2(0.001)),
    layers.Dense(32, activation="tanh", kernel_regularizer=regularizers.l2(0.001)),
    layers.Dense(1, activation="sigmoid")
])
model_three_hidden_layers.compile(optimizer="rmsprop", loss="binary_crossentropy", metrics=["accuracy"])
history_three_hidden_layers = model_three_hidden_layers.fit(partial_x_train, partial_y_train, epochs=20, batch_

# Plot the results for three hidden layers
plot_history(history_three_hidden_layers, "Three Hidden Layers")
```

Epoch 1/20
40/40 ————— **4s** 44ms/step - accuracy: 0.5028 - loss: 0.8572 - val_accuracy: 0.5016 - val_loss: 0.8123
Epoch 2/20
40/40 ————— **2s** 17ms/step - accuracy: 0.5133 - loss: 0.8084 - val_accuracy: 0.4968 - val_loss: 0.8065
Epoch 3/20
40/40 ————— **1s** 15ms/step - accuracy: 0.5141 - loss: 0.7985 - val_accuracy: 0.4928 - val_loss: 0.7976
Epoch 4/20
40/40 ————— **1s** 9ms/step - accuracy: 0.5110 - loss: 0.7922 - val_accuracy: 0.5072 - val_loss: 0.7881
Epoch 5/20
40/40 ————— **1s** 8ms/step - accuracy: 0.5164 - loss: 0.7856 - val_accuracy: 0.5084 - val_loss: 0.7837
Epoch 6/20
40/40 ————— **1s** 11ms/step - accuracy: 0.5217 - loss: 0.7788 - val_accuracy: 0.4954 - val_loss: 0.7794
Epoch 7/20
40/40 ————— **1s** 12ms/step - accuracy: 0.5234 - loss: 0.7735 - val_accuracy: 0.4942 - val_loss: 0.7749
Epoch 8/20
40/40 ————— **1s** 12ms/step - accuracy: 0.5081 - loss: 0.7690 - val_accuracy: 0.4992 - val_loss: 0.7693
Epoch 9/20
40/40 ————— **0s** 5ms/step - accuracy: 0.5203 - loss: 0.7640 - val_accuracy: 0.5038 - val_loss: 0.7624
Epoch 10/20
40/40 ————— **0s** 5ms/step - accuracy: 0.5248 - loss: 0.7582 - val_accuracy: 0.5004 - val_loss: 0.7606
Epoch 11/20
40/40 ————— **0s** 4ms/step - accuracy: 0.5215 - loss: 0.7551 - val_accuracy: 0.5106 - val_loss: 0.7540
Epoch 12/20
40/40 ————— **0s** 5ms/step - accuracy: 0.5270 - loss: 0.7505 - val_accuracy: 0.5024 - val_loss: 0.7503
Epoch 13/20
40/40 ————— **0s** 5ms/step - accuracy: 0.5241 - loss: 0.7468 - val_accuracy: 0.5066 - val_loss: 0.7471
Epoch 14/20
40/40 ————— **0s** 5ms/step - accuracy: 0.5216 - loss: 0.7433 - val_accuracy: 0.5042 - val_loss: 0.7440
Epoch 15/20
40/40 ————— **0s** 5ms/step - accuracy: 0.5219 - loss: 0.7403 - val_accuracy: 0.5052 - val_loss: 0.7412
Epoch 16/20
40/40 ————— **0s** 5ms/step - accuracy: 0.5239 - loss: 0.7379 - val_accuracy: 0.5108 - val_loss: 0.7387
Epoch 17/20
40/40 ————— **0s** 5ms/step - accuracy: 0.5255 - loss: 0.7344 - val_accuracy: 0.5054 - val_loss: 0.7355
Epoch 18/20
40/40 ————— **0s** 4ms/step - accuracy: 0.5264 - loss: 0.7322 - val_accuracy: 0.5040 - val_loss: 0.7333
Epoch 19/20
40/40 ————— **0s** 5ms/step - accuracy: 0.5195 - loss: 0.7298 - val_accuracy: 0.5060 - val_loss: 0.7301
Epoch 20/20
40/40 ————— **0s** 5ms/step - accuracy: 0.5281 - loss: 0.7270 - val_accuracy: 0.5062 - val_loss: 0.7292





Model with More Hidden Units (64)

- **Overview:** This model used 64 units per layer, increasing the model capacity.
- **Performance:** Training accuracy improved, but validation accuracy showed fluctuations. The model seemed to generalize slightly better than previous models.
- **Conclusion:** Increasing the number of units per layer showed marginal improvements, though more epochs or regularization could enhance generalization.

```
In [ ]: # Model with 64 hidden units
model_more_hidden_units = models.Sequential([
    layers.Dense(64, activation="tanh", kernel_regularizer=regularizers.l2(0.001)),
    layers.Dense(64, activation="tanh", kernel_regularizer=regularizers.l2(0.001)),
    layers.Dense(1, activation="sigmoid")
])
model_more_hidden_units.compile(optimizer="rmsprop", loss="binary_crossentropy", metrics=["accuracy"])
history_more_hidden_units = model_more_hidden_units.fit(partial_x_train, partial_y_train, epochs=20, batch_size=128)

# Plot the results for 64 hidden units
plot_history(history_more_hidden_units, "More Hidden Units (64)")
```

Epoch 1/20
40/40 2s 12ms/step - accuracy: 0.4977 - loss: 0.9042 - val_accuracy: 0.5010 - val_loss: 0.8591

Epoch 2/20
40/40 0s 7ms/step - accuracy: 0.5100 - loss: 0.8479 - val_accuracy: 0.4996 - val_loss: 0.8426

Epoch 3/20
40/40 0s 7ms/step - accuracy: 0.5172 - loss: 0.8366 - val_accuracy: 0.5056 - val_loss: 0.8369

Epoch 4/20
40/40 0s 10ms/step - accuracy: 0.5206 - loss: 0.8271 - val_accuracy: 0.4950 - val_loss: 0.8292

Epoch 5/20
40/40 1s 10ms/step - accuracy: 0.5243 - loss: 0.8177 - val_accuracy: 0.4900 - val_loss: 0.8317

Epoch 6/20
40/40 1s 10ms/step - accuracy: 0.5258 - loss: 0.8124 - val_accuracy: 0.4950 - val_loss: 0.8205

Epoch 7/20
40/40 0s 10ms/step - accuracy: 0.5226 - loss: 0.8070 - val_accuracy: 0.5124 - val_loss: 0.8091

Epoch 8/20
40/40 1s 11ms/step - accuracy: 0.5233 - loss: 0.8017 - val_accuracy: 0.5066 - val_loss: 0.8208

Epoch 9/20
40/40 0s 6ms/step - accuracy: 0.5331 - loss: 0.7942 - val_accuracy: 0.4908 - val_loss: 0.8072

Epoch 10/20
40/40 0s 5ms/step - accuracy: 0.5283 - loss: 0.7913 - val_accuracy: 0.5004 - val_loss: 0.7949

Epoch 11/20
40/40 0s 6ms/step - accuracy: 0.5271 - loss: 0.7866 - val_accuracy: 0.5038 - val_loss: 0.8044

Epoch 12/20
40/40 0s 7ms/step - accuracy: 0.5265 - loss: 0.7843 - val_accuracy: 0.4936 - val_loss: 0.8028

Epoch 13/20
40/40 0s 5ms/step - accuracy: 0.5327 - loss: 0.7777 - val_accuracy: 0.4986 - val_loss: 0.7848

Epoch 14/20
40/40 0s 7ms/step - accuracy: 0.5395 - loss: 0.7728 - val_accuracy: 0.5034 - val_loss: 0.7807

Epoch 15/20
40/40 0s 5ms/step - accuracy: 0.5411 - loss: 0.7697 - val_accuracy: 0.5010 - val_loss: 0.7793

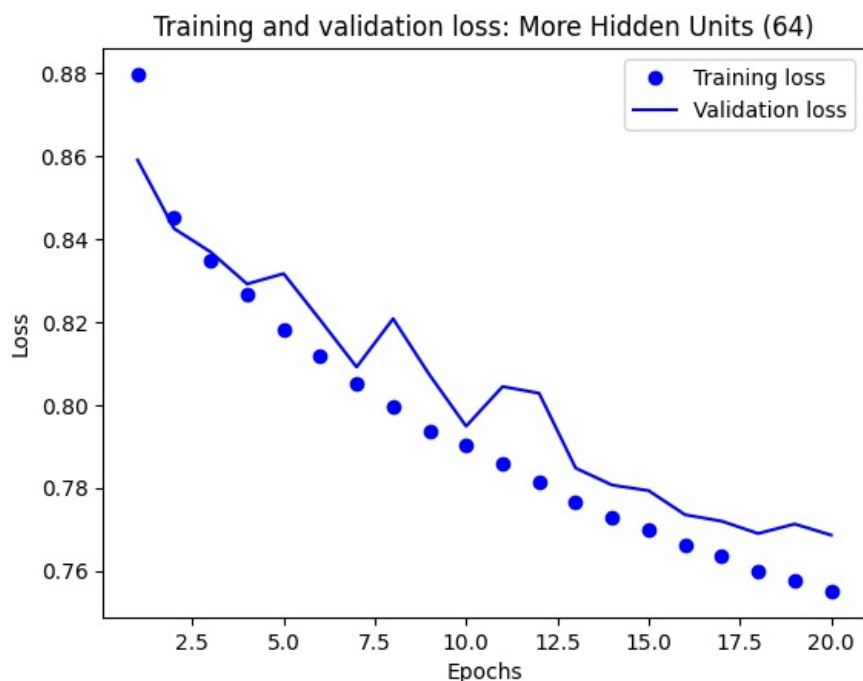
Epoch 16/20
40/40 0s 7ms/step - accuracy: 0.5452 - loss: 0.7653 - val_accuracy: 0.5076 - val_loss: 0.7735

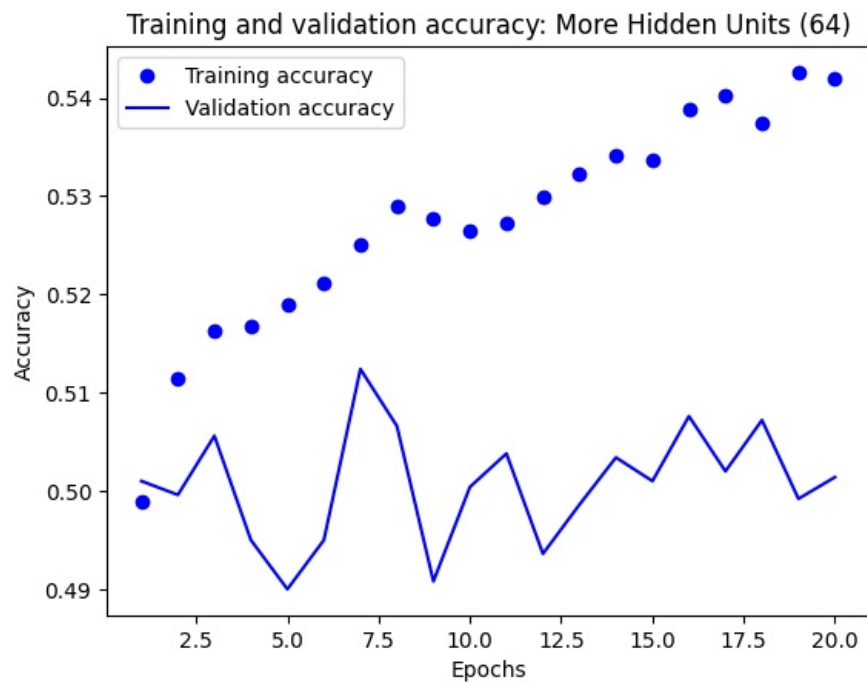
Epoch 17/20
40/40 0s 6ms/step - accuracy: 0.5448 - loss: 0.7628 - val_accuracy: 0.5020 - val_loss: 0.7720

Epoch 18/20
40/40 0s 6ms/step - accuracy: 0.5395 - loss: 0.7598 - val_accuracy: 0.5072 - val_loss: 0.7690

Epoch 19/20
40/40 0s 7ms/step - accuracy: 0.5446 - loss: 0.7577 - val_accuracy: 0.4992 - val_loss: 0.7713

Epoch 20/20
40/40 0s 7ms/step - accuracy: 0.5374 - loss: 0.7566 - val_accuracy: 0.5014 - val_loss: 0.7686





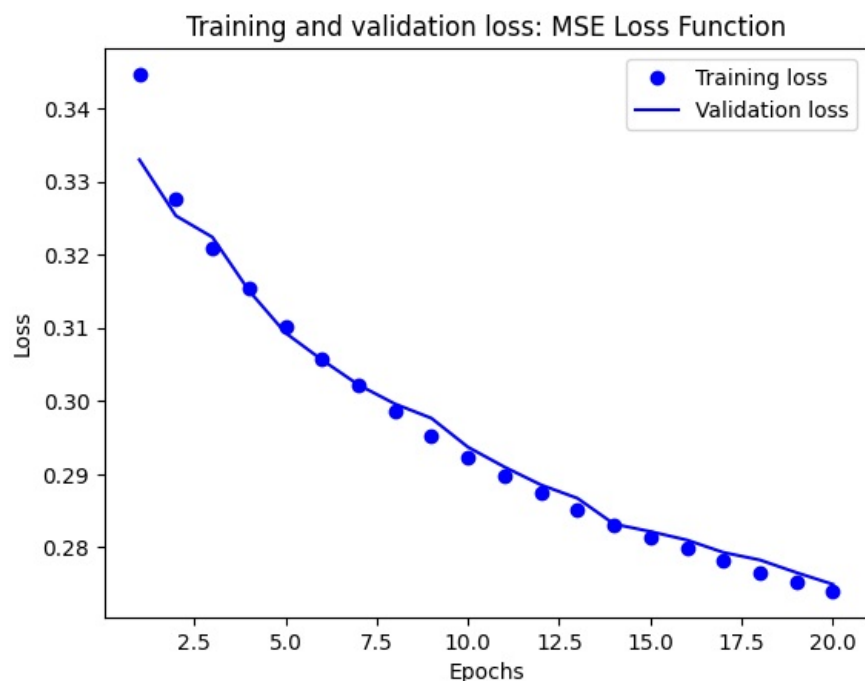
Model with MSE Loss Function

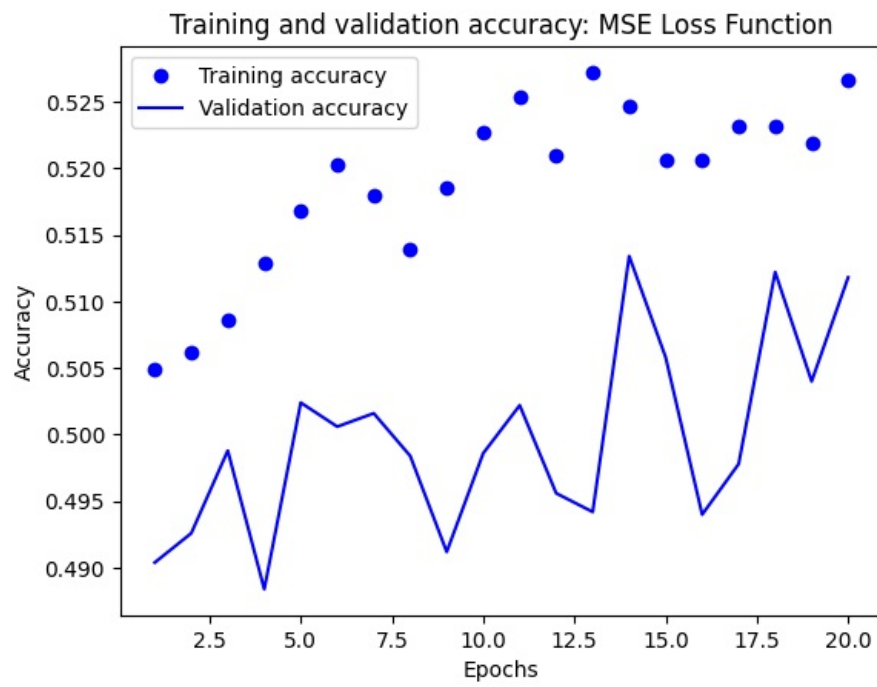
- **Overview:** This model replaced the default `binary_crossentropy` loss with `mse`.
- **Performance:** While training loss steadily decreased, validation accuracy was unstable. The model performed similarly to the `binary_crossentropy` model but did not show substantial improvements.
- **Conclusion:** MSE as a loss function did not significantly improve the results, likely because `binary_crossentropy` is more suitable for binary classification tasks.

```
In [ ]: # Model with MSE loss function
model_mse_loss = models.Sequential([
    layers.Dense(32, activation="tanh", kernel_regularizer=regularizers.l2(0.001)),
    layers.Dense(32, activation="tanh", kernel_regularizer=regularizers.l2(0.001)),
    layers.Dense(1, activation="sigmoid")
])
model_mse_loss.compile(optimizer="rmsprop", loss="mse", metrics=["accuracy"])
history_mse_loss = model_mse_loss.fit(partial_x_train, partial_y_train, epochs=20, batch_size=512, validation_d

# Plot the results for MSE loss
plot_history(history_mse_loss, "MSE Loss Function")
```

Epoch 1/20
40/40 ————— **1s** 10ms/step - accuracy: 0.5037 - loss: 0.3549 - val_accuracy: 0.4904 - val_loss: 0.3330
Epoch 2/20
40/40 ————— **0s** 4ms/step - accuracy: 0.5037 - loss: 0.3300 - val_accuracy: 0.4926 - val_loss: 0.3253
Epoch 3/20
40/40 ————— **0s** 5ms/step - accuracy: 0.5125 - loss: 0.3215 - val_accuracy: 0.4988 - val_loss: 0.3224
Epoch 4/20
40/40 ————— **0s** 4ms/step - accuracy: 0.5179 - loss: 0.3164 - val_accuracy: 0.4884 - val_loss: 0.3151
Epoch 5/20
40/40 ————— **0s** 4ms/step - accuracy: 0.5190 - loss: 0.3109 - val_accuracy: 0.5024 - val_loss: 0.3093
Epoch 6/20
40/40 ————— **0s** 4ms/step - accuracy: 0.5195 - loss: 0.3065 - val_accuracy: 0.5006 - val_loss: 0.3056
Epoch 7/20
40/40 ————— **0s** 4ms/step - accuracy: 0.5165 - loss: 0.3027 - val_accuracy: 0.5016 - val_loss: 0.3022
Epoch 8/20
40/40 ————— **0s** 5ms/step - accuracy: 0.5141 - loss: 0.2990 - val_accuracy: 0.4984 - val_loss: 0.2996
Epoch 9/20
40/40 ————— **0s** 4ms/step - accuracy: 0.5159 - loss: 0.2961 - val_accuracy: 0.4912 - val_loss: 0.2976
Epoch 10/20
40/40 ————— **0s** 5ms/step - accuracy: 0.5235 - loss: 0.2925 - val_accuracy: 0.4986 - val_loss: 0.2937
Epoch 11/20
40/40 ————— **0s** 4ms/step - accuracy: 0.5258 - loss: 0.2902 - val_accuracy: 0.5022 - val_loss: 0.2910
Epoch 12/20
40/40 ————— **0s** 4ms/step - accuracy: 0.5226 - loss: 0.2879 - val_accuracy: 0.4956 - val_loss: 0.2885
Epoch 13/20
40/40 ————— **0s** 5ms/step - accuracy: 0.5328 - loss: 0.2849 - val_accuracy: 0.4942 - val_loss: 0.2867
Epoch 14/20
40/40 ————— **0s** 4ms/step - accuracy: 0.5244 - loss: 0.2832 - val_accuracy: 0.5134 - val_loss: 0.2832
Epoch 15/20
40/40 ————— **0s** 5ms/step - accuracy: 0.5260 - loss: 0.2815 - val_accuracy: 0.5058 - val_loss: 0.2822
Epoch 16/20
40/40 ————— **0s** 5ms/step - accuracy: 0.5229 - loss: 0.2800 - val_accuracy: 0.4940 - val_loss: 0.2810
Epoch 17/20
40/40 ————— **0s** 7ms/step - accuracy: 0.5212 - loss: 0.2785 - val_accuracy: 0.4978 - val_loss: 0.2793
Epoch 18/20
40/40 ————— **0s** 7ms/step - accuracy: 0.5208 - loss: 0.2771 - val_accuracy: 0.5122 - val_loss: 0.2783
Epoch 19/20
40/40 ————— **0s** 7ms/step - accuracy: 0.5274 - loss: 0.2752 - val_accuracy: 0.5040 - val_loss: 0.2765
Epoch 20/20
40/40 ————— **1s** 15ms/step - accuracy: 0.5311 - loss: 0.2741 - val_accuracy: 0.5118 - val_loss: 0.2749





Overall Summary

In this experiment, altering the number of hidden layers, units, and loss function led to different outcomes. While increasing model complexity generally improved training performance, validation accuracy remained low, possibly indicating overfitting or insufficient regularization.